xts FAQ

xts Deveopment Team*†

July 16, 2018

Contents

What is xts?

xts is an **R** package offering a number of functionalities to work on time-indexed data. **xts** extends **zoo**, another popular package for time-series analysis.

Why should I use xts rather than zoo or another time-series package?

The main benefit of **xts** is its seamless compatibility with other packages using different time-series classes (**timeSeries**, **zoo**, ...). In addition, **xts** allows the user to add custom attributes to any object. See the main **xts** vignette for more information.

How do I install xts?

xts depends on **zoo** and suggests some other packages. You should be able to install **xts** and all the other required components by simply calling `install.packages('pkg')` from the **R** prompt.

I have multiple .csv time-series files that I need to load in a single xts object. What is the most efficient way to import the files?

If the files have the same format, load them with `read.zoo` and then call `rbind` to join the series together; finally, call `as.xts` on the result. Using a combination of `lapply` and `do.call` can accomplish this with very little code:

```
> filenames <- c("a.csv", "b.csv", "c.csv")
> sample.xts <- as.xts(do.call("rbind", lapply(filenames, read.zoo)))
```

Why is xts implemented as a matrix rather than a data frame?

xts uses a matrix rather than data.frame because:

1. **xts** is a subclass of **zoo**, and that's how **zoo** objects are structured; and

2. matrix objects have much better performance than data.frames.

How can I simplify the syntax when referring to xts object column names?

`with` allows you to use the colmn names while avoiding the full square brackets syntax. For example:

```
> lm(sample.xts[, "Res"] ~ sample.xts[, "ThisVar"] + sample.xts[, "ThatVar"])
```

can be converted to

```
> with(sample.xts, lm(Res ~ ThisVar + ThatVar))
```

How can I replace the zeros in an xts object with the last non-zero value in the series?

Convert the zeros to NA and then use na.locf:

```
> sample.xts <- xts(c(1:3, 0, 0, 0), as.POSIXct("1970-01-01")+0:5)
> sample.xts[sample.xts==0] <- NA
> cbind(orig=sample.xts, locf=na.locf(sample.xts))

                    orig locf
1970-01-01 00:00:00    1    1
1970-01-01 00:00:01    2    2
1970-01-01 00:00:02    3    3
1970-01-01 00:00:03   NA    3
1970-01-01 00:00:04   NA    3
1970-01-01 00:00:05   NA    3
```

How do I create an xts index with millisecond precision?

Milliseconds in **xts** indexes are stored as decimal values. This example builds an index spaced by 100 milliseconds, starting at the current system time:

```
> data(sample_matrix)
> sample.xts <- xts(1:10, seq(as.POSIXct("1970-01-01"), by=0.1, length=10))
```

I have a millisecond-resolution index, but the milliseconds aren't displayed. What went wrong?

Set the digits.secs option to some sub-second precision. Continuing from the previous example, if you are interested in milliseconds:

```
> options(digits.secs=3)
> head(sample.xts)

                      [,1]
1970-01-01 00:00:00.0    1
1970-01-01 00:00:00.1    2
1970-01-01 00:00:00.2    3
1970-01-01 00:00:00.3    4
1970-01-01 00:00:00.4    5
1970-01-01 00:00:00.5    6
```

I set `digits.sec=3`, but **R** doesn't show the values correctly.

Sub-second values are stored with approximately microsecond precision. Setting the precision to only 3 decimal hides the full index value in microseconds and might be tricky to interpret depending how the machine rounds the millisecond (3rd) digit. Set the `digits.secs` option to a value higher than 3 or convert the date-time to numeric and use `print`'s `digits` argument, or `sprintf` to display the full value. For example:

```
> dt <- as.POSIXct("2012-03-20 09:02:50.001")
> print(as.numeric(dt), digits=20)

[1] 1332234170.0009999275

> sprintf("%20.10f", dt)

[1] "1332234170.0009999275"
```

I am using `apply` to run a custom function on my xts object. Why does the returned matrix have different dimensions than the original one?

When working on rows, `apply` returns a transposed version of the original matrix. Simply call `t` on the returned matrix to restore the original dimensions:

```
> sample.xts.2 <- xts(t(apply(sample.xts, 1, myfun)), index(sample.xts))
```

I have an xts object with varying numbers of observations per day (e.g., one day might contain 10 observations, while another day contains 20 observations). How can I apply a function to all observations for each day?

You can use `apply.daily`, or `period.apply` more generally:

```
> sample.xts <- xts(1:50, seq(as.POSIXct("1970-01-01"),
+   as.POSIXct("1970-01-03")-1, length=50))
> apply.daily(sample.xts, mean)

                        [,1]
1970-01-01 23:30:36.244   13
1970-01-02 23:59:59.000   38

> period.apply(sample.xts, endpoints(sample.xts, "days"), mean)

                        [,1]
1970-01-01 23:30:36.244   13
1970-01-02 23:59:59.000   38

> period.apply(sample.xts, endpoints(sample.xts, "hours", 6), mean)
```

```
                         [,1]
1970-01-01 05:52:39.061  4.0
1970-01-01 11:45:18.122 10.5
1970-01-01 17:37:57.183 16.5
1970-01-01 23:30:36.244 22.5
1970-01-02 05:23:15.306 28.5
1970-01-02 11:15:54.367 34.5
1970-01-02 17:08:33.428 40.5
1970-01-02 23:59:59.000 47.0
```

How can I process daily data for a specific time subset?

First use time-of-day subsetting to extract the time range you want to work on (note the leading
"T" and leading zeros are required for each time in the range: "T06:00"), then use apply.daily
to apply your function to the subset:

```
> apply.daily(sample.xts['T06:00/T17:00',], mean)
```

How can I analyze my irregular data in regular blocks, adding observations for each regular block if one doesn't exist in the origianl time-series object?

Use align.time to round-up the indexes to the periods you are interested in, then call pe-
riod.apply to apply your function. Finally, merge the result with an empty xts object that
contains all the regular index values you want:

```
> sample.xts <- xts(1:6, as.POSIXct(c("2009-09-22 07:43:30",
+    "2009-10-01 03:50:30", "2009-10-01 08:45:00", "2009-10-01 09:48:15",
+    "2009-11-11 10:30:30", "2009-11-11 11:12:45")))
> # align index into regular (e.g. 3-hour) blocks
> aligned.xts <- align.time(sample.xts, n=60*60*3)
> # apply your function to each block
> count <- period.apply(aligned.xts, endpoints(aligned.xts, "hours", 3), length)
> # create an empty xts object with the desired regular index
> empty.xts <- xts(, seq(start(aligned.xts), end(aligned.xts), by="3 hours"))
> # merge the counts with the empty object
> head(out1 <- merge(empty.xts, count))

                    count
2009-09-22 09:00:00     1
2009-09-22 12:00:00    NA
2009-09-22 15:00:00    NA
2009-09-22 18:00:00    NA
2009-09-22 21:00:00    NA
2009-09-23 00:00:00    NA

> # or fill with zeros
> head(out2 <- merge(empty.xts, count, fill=0))
```

```
                      count
2009-09-22 09:00:00      1
2009-09-22 12:00:00      0
2009-09-22 15:00:00      0
2009-09-22 18:00:00      0
2009-09-22 21:00:00      0
2009-09-23 00:00:00      0
```

Why do I get a zoo object when I call `transform` on my xts object?

There's no **xts** method for `transform`, so the **zoo** method is dispatched. The **zoo** method explicitly creates a new **zoo** object. To convert the transformed object back to an **xts** object wrap the `transform` call in `as.xts`:

```
> sample.xts <- as.xts(transform(sample.xts, ABC=1))
```

You might also have to reset the index timezone:

```
> indexTZ(sample.xts) <- Sys.getenv("TZ")
```

Why can't I use the & operator in xts objects when querying dates?

`"2011-09-21"` is not a logical vector and cannot be coerced to a logical vector. See `?"&"` for details.

xts' ISO-8601 style subsetting is nice, but there's nothing we can do to change the behavior of `.Primitive("&")`. You can do something like this though:

```
> sample.xts[sample.xts$Symbol == "AAPL" & index(sample.xts) == as.POSIXct("2011-09-21"),]
```

or:

```
> sample.xts[sample.xts$Symbol == "AAPL"]['2011-09-21']
```

How do I subset an xts object to only include weekdays (excluding Saturday and Sundays)?

Use `.indexwday` to only include Mon-Fri days:

```
> data(sample_matrix)
> sample.xts <- as.xts(sample_matrix)
> wday.xts <- sample.xts[.indexwday(sample.xts) %in% 1:5]
> head(wday.xts)

               Open     High      Low    Close
2007-01-02 50.03978 50.11778 49.95041 50.11778
2007-01-03 50.23050 50.42188 50.23050 50.39767
```

```
2007-01-04 50.42096 50.42096 50.26414 50.33236
2007-01-05 50.37347 50.37347 50.22103 50.33459
2007-01-08 50.03555 50.10363 49.96971 49.98806
2007-01-09 49.99489 49.99489 49.80454 49.91333
```

I need to quickly convert a data.frame that contains the time-stamps in one of the columns. Using as.xts(Data) returns an error. How do I build my xts object?

The `as.xts` function assumes the date-time index is contained in the `rownames` of the object to be converted. If this is not the case, you need to use the `xts` constructor, which requires two arguments: a vector or a matrix carrying data and a vector of type `Date`, `POSIXct`, `chron`, ..., supplying the time index information. If you are certain the time-stamps are in a specific column, you can use:

```
> Data <- data.frame(timestamp=as.Date("1970-01-01"), obs=21)
> sample.xts <- xts(Data[,-1], order.by=Data[,1])
```

If you aren't certain, you need to explicitly reference the column name that contains the time-stamps:

```
> Data <- data.frame(obs=21, timestamp=as.Date("1970-01-01"))
> sample.xts <- xts(Data[,!grepl("timestamp",colnames(Data))],
+   order.by=Data$timestamp)
```

I have two time-series with different frequency. I want to combine the data into a single xts object, but the times are not exactly aligned. I want to have one row in the result for each ten minute period, with the time index showing the beginning of the time period.

`align.time` creates evenly spaced time-series from a set of indexes, `merge` ensure two time-series are combined in a single **xts** object with all original columns and indexes preserved. The new object has one entry for each timestamp from both series and missing values are replaced with `NA`.

```
> x1 <- align.time(xts(Data1$obs, Data1$timestamp), n=600)
> x2 <- align.time(xts(Data2$obs, Data2$timestamp), n=600)
> merge(x1, x2)
```

Why do I get a warning when running the code below?

```
> data(sample_matrix)
> sample.xts <- as.xts(sample_matrix)
> sample.xts["2007-01"]$Close <- sample.xts["2007-01"]$Close + 1
```

```
> #Warning message:
> #In NextMethod(.Generic) :
> #  number of items to replace is not a multiple of replacement length
```

This code creates two calls to the subset-replacement function `xts:::`[<-.xts``. The first call replaces the value of `Close` in a temporary copy of the first row of the object on the left-hand-side of the assignment, which works fine. The second call tries to replace the first *element* of the object on the left-hand-side of the assignment with the modified temporary copy of the first row. This is the problem.

For the command to work, there needs to be a comma in the first subset call on the left-hand-side:

```
> sample.xts["2007-01",]$Close <- sample.xts["2007-01"]$Close + 1
```

This isn't encouraged, because the code isn't clear. Simply remember to subset by column first, then row, if you insist on making two calls to the subset-replacement function. A cleaner and faster solution is below. It's only one function call and it avoids the `$` function (which is marginally slower on xts objects).

```
> sample.xts["2007-01","Close"] <- sample.xts["2007-01","Close"] + 1
```